

# Deep Learning Service for Efficient Data Distribution Aware Sorting

Xiaoke Zhu  
Beihang University  
China  
zhuxk@buaa.edu.cn

Qi Zhang  
Meta Platforms  
USA  
qizhang@meta.com

Wei Zhou\*  
Yunnan University  
China  
zwei@ynu.edu.cn

Ling Liu  
Georgia Institute of Technology  
USA  
ling.liu@cc.gatech.edu

**Abstract**—In this paper, we present a neural network-enabled data distribution aware sorting method, coined as NN-sort. Our approach explores the potential of developing deep learning techniques to speed up large-scale sort operations, enabling data distribution aware sorting as a deep learning service. Compared to traditional pairwise comparison-based sorting algorithms, which sort data elements by performing pairwise operations, NN-sort leverages the neural network model to learn the data distribution and uses it to map large-scale data elements into ordered ones. Our experiments demonstrate the significant advantage of using NN-sort. Measurements on both synthetic and real-world datasets show that NN-sort yields  $2.18\times$  to  $10\times$  performance improvement over traditional sorting algorithms.

## I. INTRODUCTION

Sorting is a fundamental operation in the realm of big data, powering everything from database systems [1] to bigdata analysis [2]. As the scale of data continues to grow, traditional sorting algorithms face increasing limitations in performance. While methods such as Quick Sort family [3], Merge Sort family [4], and Radix Sort family [5], [6] have long been relied upon, their comparison-based and non-comparison-based approaches are reaching a plateau in terms of scalability.

Recent research [7]–[10] has extensively explored how deep learning models can enhance the performance of traditional big data systems and algorithms. For instance, Tim Kraska et al. introduced a learned index [9] that leverages an empirical cumulative distribution function (CDF) to improve performance of traditional data structures. They also proposed a learned approach [10], [11] to improve sorting performance.

Their method, known as SageDB Sort, utilizes a learned model to generate a roughly ordered state by predicting (mapping) the positions of elements, which is then refined by a traditional sorting algorithm like Quick Sort. However, this approach has limitations. Conflicts often arise when converting the learned model’s output, such as multiple elements being mapped to the same position, leading to performance bottlenecks (see Section IV). Resolving these conflicts—especially when numerous—can be time-consuming, making it less efficient than traditional algorithms.

The question of how to design an effective deep learning-based sorting algorithm remains unanswered. Specifically, key issues include determining which type of neural network per-

forms best for sorting, understanding the complexity of neural network-based sorting, dealing with conflicts, and minimizing operations such as data comparison and movement during the sorting process.

To address these issues, this work presents NN-sort, a neural network-based sorting algorithm designed to move beyond traditional sorting paradigms. Instead of relying solely on comparisons or inherent data characteristics, NN-sort harnesses the power of neural networks to create a data distribution-aware sorting method. By training a model on historical data to predict the sorted positions of new data, NN-sort offers a novel approach that achieves efficient and scalable sorting while incorporating an effective conflict-handling mechanism.

The architecture of NN-sort consists of three distinct phases: the input phase, where data is transformed into neural network-compatible vectors; the sorting phase, where a learned neural network iteratively refines the data order; and the polish phase, where traditional methods finalize the sorting (*i.e.*, ensuring the output matches that of traditional sorting methods). This layered approach allows NN-sort to efficiently handle large datasets while minimizing the computational overhead associated with conflicts like SageDB Sort. We systematically explore the potential of NN-sort, examining its complexity, performance, and advantages over traditional algorithms. Through rigorous experiments on both synthetic and real-world datasets, we justify the effectiveness of show how NN-sort.

The contributions of this paper are summarized as follows: (a) We investigate the opportunities and challenges of enhancing traditional sorting processes by leveraging neural network-based learning approaches. (b) We develop NN-sort, a novel neural network-based sorting method that utilizes historical data to train a data distribution-aware model. This trained model performs high-performance sorting on incoming data iteratively, with an additional touch-up process to ensure the correctness of the final result. In contrast to state-of-the-art learned sorting methods, *e.g.*, SageDB Sort, NN-sort scales effectively by reducing conflicts during CDF mapping. (c) We provide a formal analysis of NN-sort’s complexity using a cost model that clarifies the intrinsic relationship between model accuracy and sorting performance. (d) We evaluate NN-sort’s performance on both synthetic and real-world datasets. Experimental results demonstrate that NN-sort achieves up to an order of magnitude speed-up in sorting time compared to

\*corresponding author

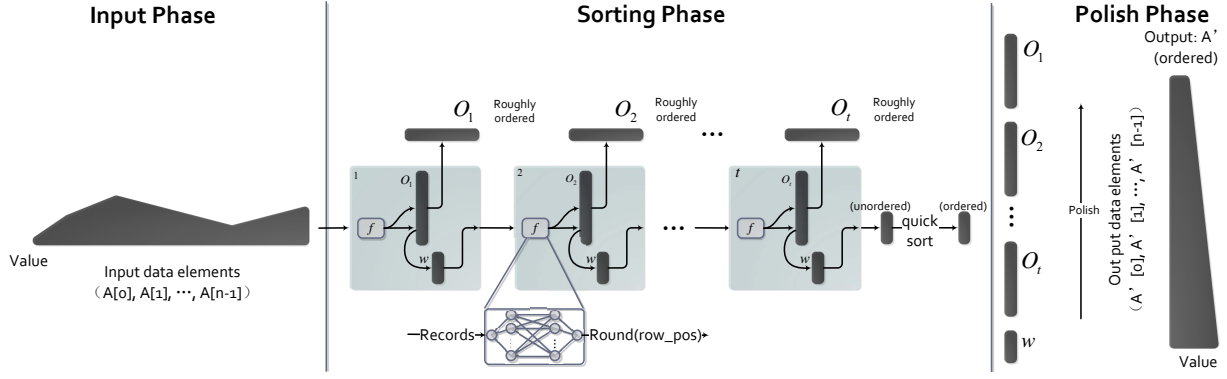


Fig. 1: NN-sort architecture

state-of-the-art sorting algorithms.

The rest of the paper is organized as follows: the design of NN-sort is presented in Section II. The time complexity and cost model of NN-sort are discussed in Section III. The experimental results are presented and analyzed in Section IV. Related work is reviewed in Section V, and the paper is concluded in Section VI.

## II. NEURAL NETWORK BASED SORT

In this section, we discuss the design of NN-sort, including how to use a neural network model for effective sorting, as well as how such a neural network model can be trained.

**Challenges.** Sorting essentially involves mapping elements from an unsorted state to a sorted state. Instead of using traditional value-comparison methods, *e.g.*, Quick Sort, this mapping can be achieved using a data distribution-aware model that takes a data element as input and produces its relative position in the sorted array as output. That is, given a function  $f$ , which returns the exact position for each key  $x \in A$ , we can sort  $A$  by determining the position of each key within the sorted order as  $\text{pos} \leftarrow f(x)$ . Before introducing the methods in detail, several challenges must be addressed to ensure this approach is effective and accurate. (1) First, for correctness, the model must precisely reflect the order among input data elements, producing results consistent with those of traditional sorting algorithms. However, it is impossible to have that magic function  $f$ , especially if we train the model just based on samples from the input data. (2) Second, for effectiveness, the model ideally should sort large volumes of data in a single pass. Achieving this requires a model to be both complex and accurate enough to capture the exact order of all elements, which can result in significant training cost and inference cost. (3) Third, conflicts arise when multiple input elements are mapped to the same output position, *i.e.*,  $f(a) = f(b)$  where  $a, b \in A$  and  $a \neq b$ . Effectively managing these conflicts is crucial for both the correctness and efficiency of the learned sorting approach. SageDB Sort addresses these collisions using traditional sorting algorithms, such as Quick Sort, which incurs computational overhead when collisions is too large. There is still room for improvement.

### Procedure NN-sort

**Input:**  $A$ : the array of data elements  
 $f$ : the learned model  
 $m$ : the relaxation factor  
 $\epsilon$  the predefined iteration limit  
 $\tau$  the predefined threshold  
**Output:**  $A'$ : the array of data elements after sorted

1.  $w \leftarrow \text{translate}(A)$
2.  $\text{init } w\_pos, O, i \leftarrow 0$
3. **while**  $0 < i < \epsilon$  and  $\text{count}(w) > \tau$  **do**
4.    $\text{init } o_i, c$
5.    $w\_pos \leftarrow w.\text{map}(f)$
6.   **for**  $\text{idx}$  in  $\text{count}(w)$  **do**
7.      $\text{pos} \leftarrow \text{round}(w\_pos[\text{idx}] * m)$
8.     **if**  $o_i[\text{pos}] == \infty$  **do**
9.        $o_i[\text{pos}] \leftarrow w[\text{idx}]$
10.    **else**
11.      $c.\text{append}(w[\text{idx}])$
12.     $O.\text{append}(o_i)$
13.     $w \leftarrow c$
14.     $++i$
15.  $\text{QuickSort}(w)$
16.  $A' \leftarrow \text{polish}(O, w)$
17. **Return**  $A'$

Fig. 2: Algorithm NN-sort

**NN-sort Design** In response to these challenges, we designed an iterative neural network-based sorting method. Unlike SageDB Sort, which employs a complex model to sort data in a single round, our approach utilizes a simpler model to sort over multiple rounds. Each round generates a roughly sorted array, with conflicts carried forward to the next iteration. This process continues until the conflicts in that iteration fall below a predefined threshold or a set number of iterations is reached. The final conflict array is then sorted using a traditional method like Quick Sort and merged with the roughly sorted arrays. After a final traversal to ensure correctness, a fully sorted result is obtained. The benefits are twofold: (1) using a

---

**Procedure** `polish( $O$ )`

*Input:*  $O = \{o_1, o_2, \dots\}$ : array of roughly sorted arrays.

$w$ : strictly sorted array.

*Output:*  $A'$ : the array of data elements after sorted

1.  $A' \leftarrow w$
  2. **for**  $o_i \in O$  **do**
  3.      $A' \leftarrow \text{InsertSort}(A', o_i)$
  4. **Return**  $A'$
- 

Fig. 3: Algorithm polish

simpler model reduces both inference and training costs; (2) the learned model can be applied repeatedly, avoiding direct sorting of conflicting elements with traditional methods.

Figure 1 illustrates this approach, where the input array  $A$  is sorted into  $A'$ . The process is divided into three phases: *Input*, *Sorting*, and *Polish*. Algorithm 2 details the workflow of NN-sort, with Line 1 addressing the input phase, Lines 2-16 covering the sorting phase, and Line 18 corresponding to the polish phase.

**Input Phase.** The input phase prepares the data for the neural network by encoding it appropriately, ensuring compatibility for processing. For example, string-type data is converted into ASCII values, transforming each character into its corresponding numeric representation. This encoding step is crucial, as it standardizes the data format and enables the neural network to interpret and process a wide variety of input types, such as integers, floating-point numbers, or categorical data, in a structured and efficient manner.

**Sorting Phase.** The sorting phase uses a learned model  $f$  to iteratively organize unordered data into roughly sorted arrays. The model maps each element to its expected position, stored at  $O_i$  where  $i$  is the iteration number, while conflicts—where multiple elements map to the same position—are stored in a temporary conflict array  $c$ . In subsequent iterations, elements in  $c$  are reprocessed by  $f$ . This process continues until either a predefined maximum number of iterations  $\epsilon$  is reached or the size of  $c$  drops below a threshold  $\tau$ , at which point it is sorted using a traditional algorithm like Quick Sort.

As shown in Algorithm 2 (line 4 - line 14), all the unsorted data elements in the  $w$  are fed into a neural-network model  $f$ , which returns a  $w\_pos[]$  array. Each element in  $w\_pos[]$  represents the expected position of corresponding elements of  $w$  within the final sorted array. It is worth mentioning that, instead of using  $w\_pos[i]$ , which is the direct output of  $f$ , we use  $\text{round}(w\_pos[i] * m)$ , a rounded value, to represent the position of  $w[i]$ . The reasons are twofold: (1) the outputs of  $f$  are decimals while the positions need to be integers. (2) with relaxation factor  $m$  the input data elements can be mapped into a larger space, thereby reducing the number of conflicts. line 11 deals with the conflicts when multiple input data elements are mapped to the same position. All conflicting data elements are stored in a conflict array  $c$  and used as input to  $f$  for the next iteration. Each iteration concludes at




Fig. 4: Example

line 14, yielding a roughly sorted array  $o_i$  and an updated conflict array  $c$ . As shown in line 3, iterations stop when the size of the conflict array falls below a threshold  $\tau$ . If the model  $f$  does not perform effectively, however, the conflict array may never shrink below  $\tau$  or may decrease too slowly, potentially resulting in higher overhead than traditional sorting algorithms. To prevent this, an additional threshold  $\epsilon$  limits the maximum number of iterations.

**Polish Phase** The polish phase refines the roughly sorted arrays  $O = \{o_1, o_2, \dots\}$  to ensure the correctness of the output. Algorithm 3 outlines this process, where the arrays in  $O$  are polished and merged with the strictly ordered array  $A'$ . The algorithm iterates over each array in  $O$ , merging them with  $A'$  one by one. Elements in  $o_i$  are either appended or inserted into the result, depending on their order relative to  $A'$ .

**Remark.** Since  $o_i$  is only roughly ordered, out-of-order elements are inserted into their correct positions in  $A'$ , ensuring NN-sort’s reliability despite potential errors from the learned model. Though insertion is costlier than appending, it is limited to out-of-order elements. As model accuracy improves, the polish phase incurs acceptable overhead. Section III discusses NN-sort’s complexity, with experimental results showing few out-of-order elements, yielding nearly linear performance.

**Example 1:** Figure 4 illustrates how NN-sort sorting.

Given thresholds  $\tau = 2$ ,  $\epsilon = 5$  and an unordered array  $A = \{32, 60, 31, 1, 81, 6, 88, 38, 3, 59, 37, 92, 91\}$ , NN-sort first checks if  $A$ ’s size is below  $\tau$ ; if so, a traditional sorting method is applied. Otherwise, learned sorting begins.

Here, NN-sort processes  $A$  in two rounds with a learned model. A conflict arises between elements 37 and 38, as  $f$  maps them to the same position, placing 37 in a conflict array  $c$ . At the end of the first iteration, elements 92, 3, 91, 32, and 59 are also in the conflict array.

After the first iteration, since  $w$ ’s size exceed  $\tau$  and the iteration count is below  $\epsilon$ , all elements in  $w$  are reprocessed by  $f$  in a second iteration, yielding a new sorted array  $o_2$  and a smaller conflict array  $c$ . Then a traditional sort (e.g., Quick Sort) is applied to  $w$ , and finally,  $o_1, o_2$ , and sorted  $w$  are merged in the polish phase, producing a fully ordered result.  $\square$

**Training.** While training time is not the focus, all our mod-

TABLE I: Notations

symbols	notations
$n$	the amount of data elements to be sorted
$\sigma_i$	collision rate per iteration
$e_i$	the number of data elements that were out-of-order in the $i$ -th iteration
$\epsilon$	the predefined limit of iterations
$t$	the number of completed iterations
$\theta$	The operations required for data to pass through $f$

els—whether shallow neural networks or simple linear/multivariate regression models—train quickly and perform well, as perfect position mapping (i.e., no conflicts or out-of-order elements) is unnecessary. The training and test data can differ; any learned order relationships help the model understand the sorting task.

### III. MODEL ANALYSIS

This section establishes the time complexity of NN-sort by analyzing key operations—moving, mapping, and comparing data elements. A cost model is introduced to highlight relationships among factors like conflict rate, model scale, iteration count, out-of-order rate, data volume, and required operations.

The total operations of NN-sort is expressed as a  $T(n, e, \sigma, t, \theta)$ , where:  $n$  is the number of data elements to be sorted,  $e = \{e_1, \dots, e_t\}$  is the set of probabilities, with  $e_i$  denoting the proportion of out-of-order elements in the  $i$ -th iteration,  $\sigma = \{\sigma_1, \dots, \sigma_t\}$  is the set of conflict rates, where  $\sigma_i$  represents the conflict rate in the  $i$ -th iteration,  $t$  is the number of iterations completed,  $\theta$  denotes the number of operations required for each data element to pass through the neural network. These basic notations are summarized in Table I.

As shown in Eq 1, the amount of operations for NN-sort to sort  $n$  ( $n > 1$ ) data elements is  $C_1 n^2 + C_2 n \log n + C_3 n$ .

$$T(n, e, \sigma, t, \theta) = \begin{cases} 1, & \text{if } n = 1 \\ C_1 n^2 + C_2 n \log n + C_3 n, & \text{if } n > 1 \end{cases} \quad (1)$$

$$C_1 = \left[ \frac{1}{2} \sum_{i=1}^t e_i (1 - \sigma_i) \left( \prod_{j=1}^{i-1} \sigma_j \right)^2 \right]$$

$$C_2 = \prod_{j=1}^t \sigma_j$$

$$C_3 = \sum_{i=1}^t \left[ \theta \sum_{j=1}^i \sigma_j + (1 - e_i)(1 - \sigma_i) \prod_{j=1}^{i-1} \sigma_j + \prod_{j=1}^i \sigma_j \right] + \left( \prod_{j=1}^t \sigma_j \right) \log \left( \prod_{j=1}^t \sigma_j \right)$$

In NN-sort, the majority of cost is spent in the Sorting and Polish phases. Let  $s(n)$  represent the time spent in the Sorting phase and  $p(n)$  represent the time spent in the Polish phase, we now formally analysis the complexity of NN-sort.

$s(n)$  consists of two kinds of operations: iteratively feed-

ing the data elements into a learned model  $f$  and sorting the array  $w$  at the last iterations using traditional sorting algorithms (e.g., QuickSort), the time complexity of which is  $n \log n$ . If  $n > 1$ , then  $\theta \sum_{j=1}^i \sigma_j n$  operations are required to feed data into model  $f$  in the  $i$ -th iteration. An additional  $(\prod_{j=1}^t \sigma_j) n \log (\prod_{j=1}^t \sigma_j) n$  operations are required to keep  $w$  order, since the size of conflicting array  $w$  updated in the last iteration is  $(\prod_{j=1}^t \sigma_j) n$ . Therefore, at the end of the algorithm, the total operations of  $s(n)$  is  $(\prod_{j=1}^t \sigma_j) n \log (\prod_{j=1}^t \sigma_j) n + \theta \sum_{i=1}^t \sum_{j=1}^i \sigma_j n$ .

$$T(n) = s(n) + p(n), \quad (n > 1)$$

$$\begin{aligned} &= \left( \prod_{j=1}^t \sigma_j \right) n \log \left( \prod_{j=1}^t \sigma_j \right) n + \theta \sum_{i=1}^t \sum_{j=1}^i \sigma_j n + p(n) \\ &= \left( \prod_{j=1}^t \sigma_j \right) n \log \left( \prod_{j=1}^t \sigma_j \right) n + \theta \sum_{i=1}^t \sum_{j=1}^i \sigma_j n \\ &\quad + \sum_{i=1}^t [e_i (1 - \sigma_i) \prod_{j=1}^{i-1} \sigma_j n \times \frac{\prod_{j=1}^{i-1} \sigma_j n}{2} \\ &\quad + (1 - e_i)(1 - \sigma_i) \prod_{j=1}^{i-1} \sigma_j n + \prod_{j=1}^i \sigma_j n] \\ &= \left[ \frac{1}{2} \sum_{i=1}^t e_i (1 - \sigma_i) \left( \prod_{j=1}^{i-1} \sigma_j \right)^2 \right] n^2 + \prod_{j=1}^t \sigma_j n \log n \\ &\quad + \left\{ \sum_{i=1}^t \left[ \theta \sum_{j=1}^i \sigma_j + (1 - e_i)(1 - \sigma_i) \prod_{j=1}^{i-1} \sigma_j \right. \right. \\ &\quad \left. \left. + \prod_{j=1}^i \sigma_j \right] + \left( \prod_{j=1}^t \sigma_j \right) \log \left( \prod_{j=1}^t \sigma_j \right) \right\} n \end{aligned}$$

$p(n)$  involves two tasks: correcting any out-of-order elements and merging the intermediate arrays (i.e.,  $o_1, \dots, o_t$  and  $w$ ). If no elements are out of order in  $o_i$ , NN-sort only needs to traverse the data once to merge them. However, in practice, out-of-order elements are almost inevitable, as the model  $f$  is unlikely to be 100% accurate.

For the ordered elements in  $o_i$ , NN-sort only requires appending it, with a time complexity of time complexity of  $O(1)$ . Therefore, in the  $i$ -th iteration, at most  $\prod_{j=1}^{i-1} \sigma_j n$  operations are required to complete the insertion, and at least one operation is needed to insert out-of-order elements. While, for an out-of-order element in the  $i$ -th merge iteration,  $\frac{\prod_{j=1}^{i-1} \sigma_j n}{2}$  operations are required to insert it into the final ordered result. Theoretically, assume that there are  $e_i (1 - \sigma_i) \prod_{j=1}^{i-1} \sigma_j n$  out-of-order elements in the  $i$ -th iteration. It takes a total of  $e_i (1 - \sigma_i) \prod_{j=1}^{i-1} \sigma_j n \times \frac{\prod_{j=1}^{i-1} \sigma_j n}{2}$  operations to process these elements. Correspondingly,  $(1 - e_i)(1 - \sigma_i) \prod_{j=1}^{i-1} \sigma_j n$  elements in  $o_i$  and  $\prod_{j=1}^i \sigma_j n$  elements in  $w$  remain ordered. Thus in the  $i$ -th merge iteration, a total of  $\prod_{j=1}^i \sigma_j n + (1 - e_i)(1 - \sigma_i) \prod_{j=1}^{i-1} \sigma_j n$  operations are required to append the ordered elements to the final result. Overall, NN-sort requires a total

of  $\sum_{i=1}^t [e_i(1 - \sigma_i) \prod_{j=1}^{i-1} \sigma_j n \times \frac{\prod_{j=1}^{i-1} \sigma_j n}{2} + (1 - e_i)(1 - \sigma_i) \prod_{j=1}^{i-1} \sigma_j n + \prod_{j=1}^i \sigma_j n]$  to sort  $n$  data elements (We show the detail in Equation 1).

#### IV. EXPERIMENTAL STUDY

Using real and synthetic data, we conducted five experiments to evaluate (1) overall sorting performance, (2) iteration impact, and (3) effects of changing data distribution.

##### A. Experimental setup

**Datasets.** The datasets used in this section are generated from the most commonly observed distributions in the real world, such as uniform distribution, normal distribution, and log-normal distribution. The models used in the experiments are trained over a subset of the testing data. The sizes of the testing dataset vary from 200MB to 500MB and each data element is 64 bits wide floating number. To verify the performance of the NN-sort under the real-world dataset. We use the QuickDraw game dataset from Google Creative Lab [12], which consists of 50,426,265 records of schema {'key-id', 'word', 'country code', 'timestamp', 'recognized', 'drawing'}. Sorting is perform on 'key-id'.

**Baselines.** We compared with five baselines (1) **Quick Sort** [13]: This algorithm divides the input dataset into two independent partitions, such that all the data elements in the first partition is smaller than those in the second partition. Then, the dataset in each partition is sorted recursively. The time complexity of **Quick Sort** can achieve  $\mathcal{O}(n \log n)$  in the best case while  $\mathcal{O}(n^2)$  in the worst case. (2) **std::sort** [14]: **std::sort** is one of the most widely used sorting algorithms from c++ standard library, and its time complexity is  $\mathcal{O}(n \log n)$  (3) **std::heap sort** [14]: **std::heap sort** is another sorting algorithm from c++ standard library, and it guarantees to perform at  $\mathcal{O}(n \log n)$  time complexity. (4) **Redis Sort** [15]: Redis Sort is a sorting method based on a data structure named *sortSet*. To sort  $M$  data elements in a *sortSet* of size  $N$ , the efficiency of **Redis Sort** is  $\mathcal{O}(N + M \log(M))$ . In addition, we also compared NN-sort with (5) **SageDB Sort** [10], [11], learding performance DNN-based sorting method. The relaxation factor  $m$  is set to 1.25 for learned sorting methods to reduce conflicts.

**Measurements.** We used sorting time and sorting rate of Equation 2 to evaluate the overall performance.

$$\text{sorting rate} = \frac{\text{elements}}{\text{time to finish sorting}} \quad (2)$$

We also used traditional sorting rate to evaluate learned-based sorting methods which is describe in Equation 3. This rate indicates how many data elements still require traditional sorting due to model inaccuracy in the learning-based approach. Ideally, a lower traditional sorting rate signifies better performance of the learning-based sorting.

$$\text{Traditional sorting rate} = \frac{\text{size}(\text{last conflicting array } w)}{\text{size of the original array } A} \quad (3)$$

TABLE II: Evaluation under real-world data

Algorithm name	Time (sec.)	Sorting Rate (elements/sec.)	The traditional sorting rate (%)
Quick Sort	10.86	4666.14	-
std::heap sort	13.46	3746.44	-
std::sort	23.71	2127.19	-
Redis Sort	63.14	798.6320	-
SageDB Sort	10.53	4790.125	9.16
NN-sort	8.47	5950.186	0.4

**Environment.** Experiments were conducted on a machine with 64GB RAM, a 2.6GHz Intel i7 processor, and a GTX1080Ti GPU with 16GB memory, running RedHat Enterprise Server 6.3. Each result reported is the median of ten runs.

**Training details.** We employed a regression model with three hidden layers, containing 2, 6, and 1 neurons, respectively. A rounding function is used to determine each element's final position. Adam [16] was the chosen optimizer. The training was conducted using a GPU and is performed offline, so training time is excluded from the runtime.

$$\text{loss}_\delta = \begin{cases} \frac{1}{2}(f(x_i) - \text{label}_i)^2, & \text{if } |f(x_i) - \text{label}_i| \leq \delta, \\ \delta|f(x_i) - \text{label}_i| - \frac{1}{2}\delta^2, & \text{otherwise} \end{cases} \quad (4)$$

In order to avoid the impact of outliers during training, the model used in experiments is trained according to the Huber loss [17] as shown in Equation 4. The batch size for training is set to 128. For all environments, we use the Adam optimizer with a learning rate of 0.001.

##### B. Experimental results.

**Exp-1: Overall Sorting Performance.** Figure 5 compares the efficiency of NN-sort with the other traditional sorting algorithms by using different datasets with increasing sizes. The total sorting time is shown in Figure 5(a) - Figure 5(c), the sorting rates are displayed in Figure 5(d) - Figure 5(f). Figure 5(g) - Figure 5(i) compare the traditional sorting rate between NN-sort and SageDB sort, which is defined as Equation 3.

The following key points highlight NN-sort's performance:

NN-sort demonstrates significant advantages over traditional sorting algorithms. As shown in Figure 5(d), its sort rate for a dataset with a lognormal distribution reaches nearly 8,300 elements per second, outperforming std::heap sort by 2.8 $\times$ , Redis Sort by 10.9 $\times$ , std::sort by 4.78 $\times$ , and Quick Sort by 218%. It also surpasses SageDB Sort by 15%. The dataset's range—its maximum and minimum values—affects NN-sort's performance. In Figure 5(h), a slight decline in sorting rate occurs due to concentrated values, which lead to more conflicts and degrade performance. Fewer records within the same range improve sorting efficiency. NN-sort performs better with uniform data distribution than with other distributions. For instance, its sorting rate reaches about 8,000 records per second with a uniform distribution, approximately




Fig. 5: Overall performance evaluation

1.3x higher than with a normal distribution, due to fewer conflicts in uniformly generated records. NN-sort consistently maintains a lower reliance on traditional sorting compared to SageDB Sort. More elements are accurately sorted by the neural model in NN-sort, reducing the need for traditional sorting, which is more time-consuming. This explains NN-sort's superior performance over SageDB Sort.

**Exp-2: Evaluation on real-word Dataset.** We used the model that is trained in previous subsections under uniformly distributed data to evaluate the performance of NN-sort over real world dataset, QuickDraw. As shown in Table II, NN-sort shows significant performance benefits over traditional sorting under real-world data. In terms of the sorting rate, NN-sort is 5950 per second, which is 2.72 times of std::sort and 7.34 times of Redis Sort. It is also 58% faster than std::heap sort. We can also observe that NN-sort outperforms SageDB Sort in terms of both the traditional sorting rate and sorting rate.

**Exp-3: Impact of the Changing Data Distribution.** As demonstrated in the previous experiments, NN-sort performs




Fig. 6: The impact of data distribution on NN-sort performance

well when the sorting data has a similar distribution to the training data. But what happens if the sorting data has a different distribution? To explore this, we trained a model using a 100MB dataset with uniformly distributed elements and then applied it to sort datasets with varying distributions.

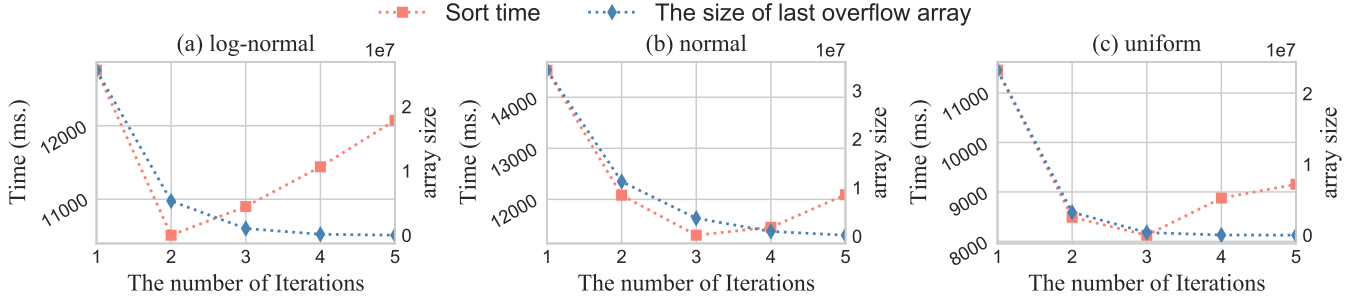


Fig. 7: Impact of Iterations

Specifically, the test dataset mixed uniformly and normally distributed data, with the latter being referred to as noisy data. We measured sorting time to evaluate the effectiveness of NN-sort, as shown in Figure 6. The results show that as the noise in the dataset increases, NN-sort’s effectiveness declines. This is due to the increasing number of out-of-order elements generated in each iteration when the distribution of the test data diverges from the training data. These elements must be sorted by traditional algorithms during the polish phase. However, NN-sort still tolerates distribution changes to some extent, outperforming the widely-used `std::sort` algorithm even with up to 45% noisy data.

**Exp-4: Impact of Iterations.** The sorting performance of NN-sort is influenced by both the size of the final conflicting array and the number of iterations. Specifically, increasing the number of iterations reduces the size of the last conflicting array that needs to be sorted using traditional methods, but it also extends the time spent on the model due to the additional iterations. Conversely, fewer iterations result in a larger conflicting array, which requires more time for traditional sorting. In this set of experiments, we quantify the impact of these factors to help users optimize NN-sort for better performance.

In Figure 7, the rhombus-dotted line represents the size of the final conflicting array, while the round-dotted line shows the total sorting time. The results indicate that although more iterations reduce the size of the final conflicting array, they don’t necessarily lead to better performance. This is because each iteration requires the model to process all input data elements. Our experiments show that 2-3 iterations strike a good balance between conflicting array size and sorting time.

**Exp-5: Evaluation of sorting accuracy.** A more complex neural network typically results in stronger model expressibility, leading to lower conflict rates, fewer out-of-order elements, but higher inference time. Therefore, it is essential to understand how these factors impact the overall time complexity of NN-sort. In this section, we use a cost curve to illustrate how the quality of the model, represented by the conflict rate  $\sigma_i$  and out-of-order rate  $e_i$  in each sorting iteration, affects the performance of NN-sort.

Figure 8 shows the number of operations required by NN-sort to sort  $n$  elements compared to Quick Sort ( $\mathcal{O}(n \log n)$ )




Fig. 8: Comparison of operations between traditional sorting algorithm and NN-sort with different model qualities.

as a baseline. We vary the model quality in NN-sort to demonstrate performance changes. The analysis assumes that NN-sort executes a maximum of five iterations ( $t = 5$ ), the model has a scale  $\theta$  of 32 neurons, and the conflict rate ( $\sigma_i$ ) equals the out-of-order rate ( $e_i$ ), remaining constant across iterations. The results indicate that NN-sort significantly outperforms Quick Sort when the conflict and out-of-order rates are 10%, and performance improves further when these rates drop to 5

In summary, fewer conflicts and misordered elements lead to more efficient sorting with NN-sort, and a well-trained model with a misorder rate of 10% or lower can outperform traditional sorting algorithms in terms of computational efficiency.

## V. RELATED WORK

Sorting is one of the most fundamental algorithms in computing. We identify two key research areas: methods to reduce sorting time complexity and neural network-based data structures.

**Methods for reducing the sorting time complexity.** Many researchers have focused on accelerating sorting by reducing its time complexity. Traditional comparison-based sorting algorithms like Quick Sort, Merge Sort, and Heap Sort require at least  $\log n! \approx n \log n - 1.44n$  operations to sort  $n$  elements [18]. Among these, Quick Sort achieves  $\mathcal{O}(n \log n)$  average complexity but degrades to  $\mathcal{O}(n^2)$  in the worst case. Merge Sort, while guaranteeing a worst-case of  $n \log n - 0.91n$ ,



requires additional linear space relative to the number of elements [18]. To mitigate the drawbacks of these algorithms and further reduce sorting time, researchers have combined different sorting techniques to leverage their strengths. Tim Sort [19], the default sorting algorithm in Java and Python, combines Merge Sort and Insertion Sort [13] to achieve fewer than  $n \log n$  comparisons on partially sorted arrays. Stefan Edelkamp et al. proposed Quickx Sort [20], which uses at most  $n \log n - 0.8358n + \mathcal{O}(\log n)$  operations for in-place sorting. They also introduced a median-of-medians variant of Quick Merge Sort [18], which employs the median-of-medians algorithm for pivot selection, reducing the operation count to  $n \log n + 1.59n + \mathcal{O}(n^{0.8})$ .

Redis Sort [15] is a build-in sorting method of the Redis database based on the sortSet data structure. It sorts  $M$  elements in a sortSet of size  $N$  with an efficiency of  $\mathcal{O}(N + M \log(M))$ .

Unlike previous work, this approach uses a learned model complexity to map an unordered array to a roughly ordered state, reducing overall operations. In the worst case, NN-sort has complexity  $\mathcal{O}(n^2)$  if all elements map to the same position, though practical operations remain lower than traditional sorting. This is validated by our experiment in Figure 8.

**Learned data structures and algorithms.** This thread of research is to explore the potential of using the neural network-based learned data structures to improve the performance of systems. Tim Kraska [9], [10] discussed the benefits of learned data structures and suggested that R-tree can be optimized by learned data structures. Xiang et al. [8] proposed an LSTM-based inverted index structure. By learning the empirical distribution function, their learned inverted index structure led to fewer average look-ups when compared with traditional inverted index structures. Alex Galakatos et al. [21] presented a data-aware index structure called FITing-Tree, which can approximate an index using piece-wise linear functions with a bounded error specified at construction time. Michael Mitzenmacher [22] proposed a learned sandwiching bloom filter structure, while the learned model is sensitive to data distributions.

Unlike the research mentioned above, our approach integrates sorting with learning by training a model to enhance sorting performance. Additionally, we employ an iteration-based mechanism to further optimize performance by minimizing conflicts. We also provide a formal analysis of the time complexity of our approach and present a cost model to balance model accuracy with sorting performance. A closely related work is SageDB Sort [11], [21], which leverages deep neural networks for sorting. Our approach improves upon SageDB Sort by offering a more efficient solution for handling position conflicts generated by the learned model.

## VI. CONCLUSION

Sorting is fundamental in big data processing. We introduce NN-sort, a neural network-based sorting method that uses historical data to sort new data, iteratively reducing sorting

conflicts—a key bottleneck in learned sorting. Our analysis includes complexity, a cost model, and the balance between model accuracy and performance. Experiments show NN-sort outperforms traditional algorithms. Future work includes enhancing NN-sort’s adaptability to changing data distributions.

## REFERENCES

- [1] G. Graefe, “Implementing sorting in database systems,” *ACM Comput. Surv.*, vol. 38, no. 3, p. 10, 2006.
- [2] R. Hilker, C. Sickinger, C. N. Pedersen, and J. Stoye, “UniMoG—a unifying framework for genomic distance calculation and sorting based on DCJ,” *Bioinformatics*, vol. 28, no. 19, pp. 2509–2511, 2012.
- [3] D. Cederman and P. Tsigas, “A practical quicksort algorithm for graphics processors,” in *Algorithms - ESA 2008*, D. Halperin and K. Mehlhorn, Eds., 2008, pp. 246–258.
- [4] A. Andersson, T. Hagerup, S. Nilsson, and R. Raman, “Sorting in linear time?” *Journal of Computer and System Sciences*, vol. 57, no. 1, pp. 74–93, 1998.
- [5] S. Bandyopadhyay and S. Sahni, “GRS - GPU radix sort for multifield records,” in *HiPC*, 2010, pp. 1–10.
- [6] J. Tang and X. Zhou, “Cardinality sorting and its bit-based operation-based optimization (in chinese),” *JOURNAL OF NANJINGUNIVERSITY OF TECHNOLOGY*, vol. 20, 2006.
- [7] X. Zhu, Q. Zhang, T. Cheng, L. Liu, W. Zhou, and J. He, “Dlb: deep learning based load balancing,” in *CLOUD*, 2021.
- [8] W. Xiang, H. Zhang, R. Cui, X. Chu, K. Li, and W. Zhou, “Pavo: A rnn-based learned inverted index, supervised or unsupervised?” *IEEE Access*, vol. 7, pp. 293–303, 2019.
- [9] T. Kraska, A. Beutel, E. H. Chi, J. Dean, and N. Polyzotis, “The case for learned index structures,” in *SIGMOD*, 2018, pp. 489–504.
- [10] T. Kraska, M. Alizadeh, A. Beutel, E. H. Chi, A. Kristo, G. Leclerc, S. Madden, H. Mao, and V. Nathan, “Sagedb: A learned database system,” in *CIDR*, 2019.
- [11] J. Ding, R. Marcus, A. Kipf, V. Nathan, A. Nrusimha, K. Vaidya, A. van Renen, and T. Kraska, “Sagedb: An instance-optimized data analytics system,” *PVLDB*, vol. 15, no. 13, 2022.
- [12] Google, “Google creative lab,” Available: <https://github.com/googlecreativelab>, google Creative Lab [Online].
- [13] T. H. Cormen, *Introduction to Algorithms, 3rd Edition*. Press.
- [14] “C++ resources network,” <http://www.cplusplus.com/>, general information about the C++ programming language, including non-technical documents and descriptions.
- [15] “Redis,” <https://redis.io/>, redis is an open source (BSD licensed), in-memory data structure store, used as a database, cache and message broker.
- [16] D. P. Kingma and J. Ba, “Adam: A method for stochastic optimization,” in *ICLR*, 2015.
- [17] P. J. Huber, “Robust estimation of a location parameter,” *Annals of Mathematical Statistics*, vol. 35, no. 1, pp. 73–101, 1964.
- [18] S. Edelkamp and A. Weiß, “Worst-case efficient sorting with quickmergesort,” in *ALENEX*, 2019, pp. 1–14.
- [19] “Python resources network,” <https://www.python.org/>, general information about the Python programming language, including non-technical documents and descriptions.
- [20] S. Edelkamp and A. Weiß, “Quickxsort: Efficient sorting with  $n \log n - 1.399n + o(n)$  comparisons on average,” in *International Computer Science Symposium in Russia*, 2014, pp. 139–152.
- [21] A. Galakatos, M. Markovitch, C. Binnig, R. Fonseca, and T. Kraska, “Fiting-tree: A data-aware index structure,” in *SIGMOD*, 2019.
- [22] M. Mitzenmacher, “A model for learned bloom filters, and optimizing by sandwiching,” *CoRR*, vol. abs/1901.00902, 2019. [Online]. Available: <http://arxiv.org/abs/1901.00902>